

OpenMP

OPEN SPECIFICATIONS FOR MULTI PROCESSING


Raghu.H.V,
Scientist -D,
C-DAC, Bengaluru




OUTLINE

- ❖ Introduction
- ❖ OpenMP Programming Model
- ❖ OpenMP Directives
- ❖ Synchronization Constructs
- ❖ Runtime Libraries
- ❖ Environment Variables

INTRODUCTION



Older processors have only one CPU core to execute instructions





Modern processors have many CPU cores to execute instructions

WHY OPENMP?

When you run a Sequential Program:

- Instructions executed on 1 core
- Other cores are idle
- Wastage of available resources – We want all cores to be used – How?

Use “OpenMP”





INTRODUCTION TO OPENMP

- Standard API for writing shared memory parallel applications in C, C++, and Fortran
- OpenMP API Consists of :
 - Compiler Directives
 - Runtime Library Routines
 - Environment variables
- Specification maintained by the OpenMP Architecture Review Board (<http://www.openmp.org>)
- Scenarios
 - Creating new program
 - Parallelizing existing one
- Use Explicit Parallelization – OpenMP
 - When compiler cannot find parallelism



CHARACTERISTICS OF OPENMP EXECUTION MODEL


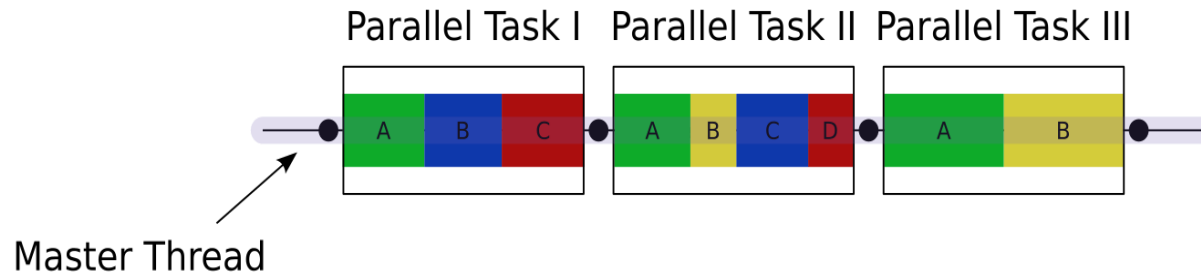
- Thread Based Parallelism
- Explicit Parallelism
- Fork - Join Model
- Compiler Directive Based



THREADS V/S PROCESS


- A process is a program in execution. A thread is a light weight process.
- Threads share the address space of the process that created it, process have their own address space.
- Threads can directly communicate with other threads of the same process. Processes must use IPC to communicate with other process.
- Changes to main thread may affect behaviour of other threads of the process. Changes to the parent process do not affect child process.

FORK-JOIN MODEL



MEMORY MODEL

- Shared Memory.
- Data is **private** or **shared**.
- Private accessed only by **owned** threads.
- **Synchronization** takes place.





HOW DO THREADS INTERACT?

- OpenMP is an multi-threading, shared memory model.
 - Threads communicate by sharing variables.
- Unintended sharing of data causes race conditions:
 - Race condition: when the program's outcome changes as the threads are scheduled differently.
- To control race conditions:
 - Use synchronization to protect data conflicts.
- Synchronization is expensive so:
 - Change how data is accessed to minimize the need for synchronization.

GENERAL SYNTAX

- Header file
 - #include <omp.h>
- Parallel Region

```
#pragma omp parallel [clauses...]  
{  
    // ... do some work here  
} // end of parallel region/block
```

- **Functions and Environment variables**

AN EXAMPLE – HELLO WORLD

Sequential

```
void main()  
{  
    int ID = 0;  
    printf(" Thread: %d - Hello World ", ID);  
}
```

OpenMP include file

OpenMP

`#include <omp.h>`

`void main()`

`{`

`#pragma omp parallel`

`{`

`int ID = 0;`

`printf(" Thread: %d - Hello World ", ID);`

`}`

End of parallel region

`}`

Parallel region with default
Number of threads



COMPILATION

- GNU Compiler Example :
 - `gcc -o helloc.x -fopenmp hello.c`
- IBM AIX compiler :
 - `xlc -o helloc.x -qsmp=omp hello.c`
- Portland group compiler:
 - `pgcc -o helloc.x -mp hello.c`
- Intel Compiler Example:
 - `icc -o helloc.x -openmp hello.c`

OPENMP COMPONENTS

Compiler Directives

- A. Parallel construct**
- B. Work Sharing**
- C. Synchronization**
- D. Data Scope**
 - **Private**
 - **shared**

Runtime Library Routines

- **Number of threads**
- **Thread ID**

Environment Variables

- **Number of threads**
- **Scheduling Type**



OPENMP COMPILER DIRECTIVES

- Used to divide blocks of code among threads.
- Distributing loop iterations among threads.
- Serializing sections of code.
- Synchronization of work among threads.



RUNTIME LIBRARY ROUTINES

- Setting and querying the number of threads.
- Querying unique thread ID.
- To check whether inside parallel region.

ENVIRONMENT VARIABLES

- Setting the number of threads.
- Specifying how loop iterations are divided.



1. OPENMP COMPILER DIRECTIVES

- Syntax:
 - `#pragma omp directive-name [clause1 clause2..] new-line`
- Example:
 - `#pragma omp parallel private(pi)`
- Example Directives
 - parallel
 - for/do
 - sections
 - single
 - critical
 - barrier



A. PARALLEL DIRECTIVE

- The parallel construct **forms a team of threads** and starts parallel execution of a parallel region.
- A parallel region is a block of code that will be executed by multiple threads.

```
#pragma omp parallel [clause1 clause2 ...] newline  
structured_block
```

Clauses:

if (scalar_expression)

private (list)

firstprivate (list)

num_threads (integer-expression)



Example: parallel construct

```
#include <omp.h>
main ()
{
    int nthreads, tid;
    /* Fork a team of threads with each thread having a private tid */
    #pragma omp parallel private(tid)
    {
        /* Obtain and print thread id */
        tid = omp_get_thread_num();
        printf("Hello World from thread = %d\n", tid);

        /* Only master thread does this */
        if (tid == 0)
        {
            nthreads = omp_get_num_threads();
            printf("Number of threads = %d\n", nthreads);
        }

    } /* Implicit Barrier - All threads join master thread and terminate */
}
```



- Main thread **creates a team of threads** and becomes the master of the team.
- The **master** is a member of that team and has **thread id 0** within that team.
- There is an **implied barrier** at the end of a parallel section.
- If any thread terminates within a parallel region, **all threads** in the team **terminate**.



DIFFERENT WAYS OF THREAD CREATION

- The number of threads in a parallel region is determined by the following factors, **in order of precedence (Low-High)**:
 1. Default number of Threads (number of cores)
 2. Setting of the **OMP_NUM_THREADS** environment variable
export OMP_NUM_THREADS=2
 3. Use of the **omp_set_num_threads()** library function
omp_set_num_threads(4);
 4. Setting of the **NUM_THREADS** clause
#pragma omp parallel private(tid) num_threads(6)
- Implementation default - usually the **number of cores**.
- Threads are numbered from 0 (master thread) to N-1






B. WORK SHARING CONSTRUCTS

- Divides execution of code region among members of the team of threads.
- Work-sharing constructs do not launch new threads
- Restrictions
 - Must be enclosed within a parallel region.
 - Work is distributed among the threads

Examples:


- for
- section
- single – serializes a portion of code



FOR DIRECTIVE

- **for directive** divide the loop iterations among threads and execute in parallel.
- **Syntax:**

```
#pragma omp for [clause1 clause2 ...] newline  
for_loop
```



Example: for construct

```
#include <omp.h>
#define N 1000

main ()
{
  int i;
  float a[N], b[N], c[N];

  /* Some initializations */
  for (i=0; i < N; i++) a[i] = b[i] = i * 1.0;

  #pragma omp parallel shared(a,b,c) private(i)
  {
    #pragma omp for
    for (i=0; i < N; i++)
      c[i] = a[i] + b[i];
  } /* end of parallel section */
}
```



RESTRICTIONS


for (index = start ; index < end ; increment_expr)

It should be possible to **determine** the number of loop **iterations** before execution, to divide the iterations.

- increment must be same at each iteration
- all loop iterations must be done.


SCHEDULE: Describes how iterations of the loop are divided among the threads in the team. The default schedule is implementation dependent.

- **STATIC**



SCHEDULE(STATIC,6)
26 iter on 4 threads

- **DYNAMIC**




SCHEDULE(DYNAMIC,6)
26 iter on 4 threads

- To set at runtime – determined by the environment variable **OMP_SCHEDULE**

SECTION DIRECTIVE

- It specifies that the enclosed section(s) of code are to be divided among the threads in the team.
- Each **SECTION** is executed only once by a thread in the team.
- Syntax:

```
#pragma omp sections [clause1 clause2]
{
    #pragma omp section <newline>
    structured_block
    #pragma omp section <newline>
    structured_block
}
```



○ Key points

- Implicit barrier at the end of sections directive, if nowait clause is present.
- If “too many” sections, some threads execute more than one section (round-robin).
- If “too few” sections, some threads are idle.
- We don’t know in advance which thread will execute which section.

```
#pragma omp parallel default(none)\  
    shared(n,a,b,c,d) private(i)
```

```
{
```

```
#pragma omp sections nowait
```

```
{
```

```
#pragma omp section
```

```
    for (i=0; i<n-1; i++)  
        b[i] = (a[i] + a[i+1])/2;
```

```
#pragma omp section
```

```
    for (i=0; i<n; i++)  
        d[i] = 1.0/c[i];
```

```
} /*-- End of sections --*/
```


```
} /*-- End of parallel region --*/
```

No need to
wait to enter
Or to exit
from a
section

SINGLE DIRECTIVE

- The enclosed code is to be **executed by only one thread** in the team.
- A **barrier is implicitly set at the end** of the single block (the barrier can be removed by the **nowait** clause)
- Syntax:

```
#pragma omp single [clause1 clause2 ...]  
structured_block
```





SYNCHRONIZATION

- Consider a case of 2 threads, both trying to increment a variable “x” at the same time (assume x = 0 initially)
- Problem: One possible execution sequence:
 - Thread 1 loads the value of x into register A.
 - Thread 2 loads the value of x into register A.
 - Thread 1 adds 1 to register A $\rightarrow x = x + 1$ by thread1
 - Thread 2 adds 1 to register A $\rightarrow x = x + 1$ by thread2
 - Thread 1 stores register A at location x
 - Thread 2 stores register A at location x
 - The resultant value of x will be 2, not 1 as it should be.



SYNCHRONIZATION

- Solution: To avoid a situation like this, the incrementing of x must be synchronized between the two threads to ensure that the correct result is produced.
- OpenMP Synchronization Constructs:
 - MASTER Directive
 - CRITICAL Directive
 - BARRIER Directive
 - ATOMIC Directive



MASTER

- To execute a region only by master thread of the team.
- All other threads on the team skip this section of code
- There is no implied barrier associated with this directive
- Syntax:
`#pragma omp master newline
structured_block`


CRITICAL

- The **CRITICAL** directive specifies a region of code that must be **executed by only one thread at a time**.

- Syntax:

```
#pragma omp critical [ name ] newline  
structured_block
```

The optional name enables multiple different CRITICAL regions to exist.





CRITICAL

- If a thread is currently executing inside a CRITICAL region and another thread reaches that CRITICAL region and attempts to execute it, it will block until the first thread exits that CRITICAL region.
- The optional name enables multiple different CRITICAL regions to exist:
 - Names act as global identifiers. Different CRITICAL regions with the same name are treated as the same region.
 - All CRITICAL sections which are unnamed, are treated as the same section.

BARRIER

- On reaching BARRIER directive, a thread will wait at that point until all other threads have reached that barrier.
- All threads then resume executing in parallel the code that follows the barrier. Synchronizes all threads in a team.
- Syntax:

`#pragma omp barrier` newline

```
#pragma omp parallel for  
for (i=0; i < N; i++)  
  
    a[i] = b[i] + c[i];
```

```
#pragma omp barrier
```

```
#pragma omp parallel for  
for (i=0; i < M; i++)  
  
    d[i] = a[i] + b[i];
```




ATOMIC

- Specifies that a specific **memory location** must be **updated** atomically, by a single thread
- **Atomicity:** We cannot split an operation – Example, A write operation.
- The directive applies only to a single, immediately following statement → **Applies to only a single statement, allows only a limited number of expressions.**
- Provides a mini critical section
- **Syntax:**
`#pragma omp atomic newline
statement_expression`

```
#pragma omp parallel shared(sum,x,y)
...
#pragma omp critical
{
    update(x);
    update(y);
    sum=sum+1;
}
...
!$OMP END PARALLEL
```

```
#pragma omp parallel shared(sum)
...
{
    #pragma omp atomic
    sum=sum+1;
    ...
}
```



- This Directive is very similar to the CRITICAL directive.
- Difference is that ATOMIC is only used for the update of a memory location.
- Sometimes ATOMIC is also referred to as a mini critical section.



DATA SCOPE ATTRIBUTES

The OpenMP **Data Scope Attribute Clauses** are used to explicitly define how **variables** should be **scoped**. They include:

- **IF**
- **PRIVATE**
- **FIRSTPRIVATE**
- **SHARED**

Data Scope Attribute Clauses are used in conjunction with several directives (PARALLEL, DO/for, and SECTIONS) to control the scoping of enclosed variables



IF

- A clause to decide at run time if a parallel region should actually be executed in parallel (multiple threads) or just by the master thread:
- Syntax:
 - If(logical expr)
- Example:
 - `#pragma omp parallel if (n>100000)`

PRIVATE

- This declares variables in its list to be private to each thread
- Syntax:
 - `private (list)`
Eg: `int B = 10;`
 `#pragma omp parallel private(B)`
 `B=...;`
- A `un-initialised copy` of B is created before the parallel region begins.



FIRSTPRIVATE

- **Firstprivate:** A private initialized copy of B is created on each thread's stack before the parallel region begins

- Syntax:

`firstprivate (list)`

- Example:

```
int B = 10;
```

```
#pragma omp parallel firstprivate(B)
```

```
B = 10;
```

- SHARED Clause

- A shared variable **exists in only one memory location** and all threads can read or write to that address
- Syntax:

shared (list)

- It is the programmer's responsibility to ensure that multiple threads properly access SHARED variables (such as via CRITICAL sections)

2. RUNTIME LIBRARIES

- **Execution environment routines** that can be used to control and to query the parallel execution environment

Example routines:

- **OMP_SET_NUM_THREADS**
 - omp_set_num_threads routine affects the number of threads to be used for subsequent parallel regions
 - C/C++ : `void omp_set_num_threads(int num_threads);`
- **OMP_GET_NUM_THREADS**
 - returns the number of threads in the current team.
 - C/C++ : `int omp_get_num_threads(void);`

○ OMP_GET_THREAD_NUM

- Returns the thread ID of the thread

```
#include <omp.h>
```

```
int omp_get_thread_num()
```

○ OMP_GET_NUM_PROCS

- To get the number of processors

```
#include <omp.h>
```

```
int omp_get_num_procs()
```

○ OMP_IN_PARALLEL

- Determine if the section of code which is executing is parallel or not.

```
#include <omp.h>
```

```
int omp_in_parallel()
```

3. ENVIRONMENT VARIABLES

- OMP_SCHEDULE
 - export OMP_SCHEDULE "static"
 - export OMP_SCHEDULE "dynamic"

- OMP_NUM_THREADS
 - export OMP_NUM_THREADS=8



Thank You